

Generative alternatives

Skorik Sergey

MIPT, 2023

October 14, 2023

1 Backgrounds

2 HiPPO

3 S4


4 Kalman Filter

Backgrounds

State-space representation

The most general state-space representation of a linear system with p inputs, q outputs and n state variables is written in the following form:

$$\begin{cases} \dot{\mathbf{x}}(t) = \mathbf{A}(t)\mathbf{x}(t) + \mathbf{B}(t)u(t) \\ \mathbf{y}(t) = \mathbf{C}(t)\mathbf{x}(t) + \mathbf{D}(t)u(t) \end{cases} \quad (1)$$



Backgrounds

State-Space representation

State-Space representation is a mathematical model of a physical system. For example, for mass spring damper system (<https://cookierobotics.com/008/>) the following equations describes them:

$$\begin{cases} \begin{bmatrix} \dot{x}_1(t) \\ \dot{x}_2(t) \end{bmatrix} = \begin{bmatrix} 0 & 1 \\ -\frac{k}{m} & -\frac{b}{m} \end{bmatrix} \cdot \begin{bmatrix} x_1(t) \\ x_2(t) \end{bmatrix} + \begin{bmatrix} 0 \\ \frac{1}{m} \end{bmatrix} \cdot [u(t)] \\ [y(t)] = \begin{bmatrix} 1 & 0 \end{bmatrix} \cdot \begin{bmatrix} x_1(t) \\ x_2(t) \end{bmatrix} \end{cases} \quad (2)$$

real-world dynamics might not be fully known or may be subject to uncertainty. In such cases, A is used to capture the probabilistic or stochastic nature of the state dynamics. Instead of being a fixed matrix, A is modeled as a distribution over possible state transitions $A \sim \mathcal{N}(\mu, \Sigma)$.

HiPPO Framework

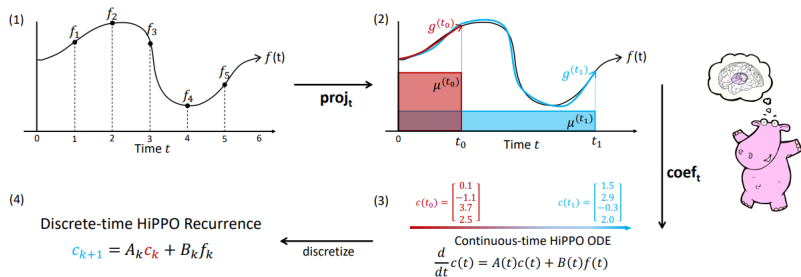


Figure 1: **Illustration of the HiPPO framework.** (1) For any function f , (2) at every time t there is an optimal projection $g^{(t)}$ of f onto the space of polynomials, with respect to a measure $\mu^{(t)}$ weighing the past. (3) For an appropriately chosen basis, the corresponding coefficients $c(t) \in \mathbb{R}^N$ representing a compression of the history of f satisfy linear dynamics. (4) Discretizing the dynamics yields an efficient closed-form recurrence for online compression of time series $(f_k)_{k \in \mathbb{N}}$.

HiPPO Framework

Connection between SSR and HiPPO

The composition $\text{coef} \circ \text{proj}$ is called hippo, which is an operator mapping a function f to the optimal projection coefficients c , i.e.

$(\text{hippo}(f))(t) = \text{coef}_t(\text{proj}_t(f))$. The coefficient function $c(t) = \text{coef}_t(\text{proj}_t(f))$ has the form of an ODE satisfying $\dot{c}(t) = A(t)c(t) + B(t)f(t)$

SSM setup

SSM setup

The state space model maps a 1-D input signal $u(t)$ to an N-D latent state $x(t)$ before projecting to a 1-D output signal $y(t)$

$$\begin{cases} \dot{x}(t) = \mathbf{A}x(t) + \mathbf{B}u(t) \\ y(t) = \mathbf{C}x(t) \end{cases} \quad (3)$$

The discrete SSM is

$$\begin{cases} x_k = \bar{\mathbf{A}}x_{k-1} + \bar{\mathbf{B}}u_k, & \bar{\mathbf{A}} = (\mathbf{I} - \Delta/2 \cdot \mathbf{A})^{-1}(\mathbf{I} + \Delta/2 \cdot \mathbf{A}) \\ y_k = \bar{\mathbf{C}}x_k, & \bar{\mathbf{B}} = (\mathbf{I} - \Delta/2 \cdot \mathbf{A})^{-1}\Delta\mathbf{B}, \quad \bar{\mathbf{C}} = \mathbf{C} \end{cases} \quad (4)$$

The fundamental bottleneck in computing the discrete-time SSM (4) is that it involves repeated matrix multiplication by $\bar{\mathbf{A}}$. For example, naively as in the LSSL involves L successive multiplications by $\bar{\mathbf{A}}$ requiring $\mathcal{O}(N^2 \cdot L)$ operations and $\mathcal{O}(NL)$ space.

HiPPO

HiPPO specifies a class of certain matrices $\mathbf{A} \in \mathbb{R}^{N \times N}$ that when incorporated into (3), allows the state $x(t)$ to memorize the history of the input $u(t)$.

$$\text{(HiPPO Matrix)} \quad \mathbf{A}_{nk} = - \begin{cases} (2n+1)^{1/2}(2k+1)^{1/2} & \text{if } n > k \\ n+1 & \text{if } n = k \\ 0 & \text{if } n < k \end{cases} \quad (5)$$

The ideal scenario is when the matrix \mathbf{A} is diagonalizable by a perfectly conditioned (i.e., unitary) matrix. By the Spectral Theorem of linear algebra, this is exactly the class **normal matrices**. In particular, it does not contain the HiPPO matrix.

Table 1: Complexity of various sequence models in terms of sequence length (L), batch size (B), and hidden dimension (H); tildes denote log factors. Metrics are parameter count, training computation, training space requirement, training parallelizability, and inference computation (for 1 sample and time-step). For simplicity, the state size N of S4 is tied to H . Bold denotes model is theoretically best for that metric. Convolutions are efficient for training while recurrence is efficient for inference, while SSMs combine the strengths of both.

	Convolution ³	Recurrence	Attention	S4
Parameters	LH	H^2	H^2	H^2
Training	$\tilde{L}H(B + H)$	BLH^2	$B(L^2H + LH^2)$	$BH(\tilde{H} + \tilde{L}) + B\tilde{L}H$
Space	BLH	BLH	$B(L^2 + HL)$	BLH
Parallel	Yes	No	Yes	Yes
Inference	LH^2	H^2	$L^2H + H^2L$	H^2

Kalman Filter

The more general formulation of the state space model described in the previous section as an observation equation

$$y_t = A_t x_t + V_t$$

and a state equation

$$x_t = \Theta x_{t-1} + W_t$$

where y_t is a $p \times 1$ vector, x_t is a $k \times 1$ vector, A_t is a $p \times k$ matrix and Θ is $k \times k$ matrix. We can think of $V_t \sim \mathcal{N}(0, S)$ and $W_t \sim \mathcal{N}(0, R)$.

Given an initial state x_0^0 and P_0^0 , the prediction equations are (analogous to above)

$$\begin{aligned} x_1^0 &= \Theta x_0^0 \\ P_1^0 &= \Theta P_0^0 \Theta' + R \end{aligned}$$

and the updating equations are, given a new observation y_1 ,

$$\begin{aligned} x_1^1 &= x_1^0 + K_1(y_1 - A_1 x_1^0) \\ P_1^1 &= (I - K_1 A_1) P_1^0 \end{aligned}$$

where

$$K_1 = P_1^0 A_1' (A_1 P_1^0 A_1' + S)^{-1}.$$

In general, given the current state x_{t-1}^{t-1} and P_{t-1}^{t-1} and a new observation y_t , we have

$$\begin{aligned} x_t^{t-1} &= \Theta x_{t-1}^{t-1} \\ P_t^{t-1} &= \Theta P_{t-1}^{t-1} \Theta' + R \end{aligned}$$